

```
[9]: import pandas as pd
import seaborn as sns
import numpy as np
import statsmodels.api as sm
import matplotlib.pyplot as plt
from statsmodels.stats.outliers_influence import variance_inflation_factor
from scipy import stats
```

```
[10]: print("""Answer to question #1""")
data = {
    'Dependent': [35, 50, 65, 70, 80],
    'education': [12, 16, 18, 20, 21],
    'experience': [5, 10, 12, 15, 18],
    'Age': [25, 30, 32, 35, 40]
}

df = pd.DataFrame(data)
print(df)

Y=df ["Dependent"]
X=df[["education", "experience", "Age"]]
X=sm.add_constant(X)
model = sm.OLS(Y, X).fit()
print(model.summary())
print("-----")
print("-----")
print(f"\nInterpretation:")
print("""In the abstract sense, the coefficient of x1 that is, the coefficient
on education indicates that for every additional unit of education, the
dependent variable y increases by approximately 15.8333 on avg""")
```

Answer to question #1:

	Dependent	education	experience	Age
0	35	12	5	25
1	50	16	10	30
2	65	18	12	32
3	70	20	15	35
4	80	21	18	40

```

OLS Regression Results
=====
Dep. Variable:                 Dependent   R-squared:                   0.997
Model:                          OLS         Adj. R-squared:             0.987
Method:                         Least Squares   F-statistic:                  99.67
Date:                          Sun, 02 Nov 2025   Prob (F-statistic):        0.0735
Time:                           20:10:55     Log-Likelihood:            -6.6389
No. Observations:                  5         AIC:                      21.28
Df Residuals:                      1         BIC:                      19.72
Df Model:                        3
Covariance Type:                nonrobust
=====
              coef    std err          t      P>|t|      [0.025      0.975]
-----
const      -335.0000    147.549     -2.270      0.264    -2209.793    1539.793
education     15.8333     6.067      2.610      0.233     -61.252     92.919
experience    -20.4167    9.887     -2.065      0.287    -146.037    105.203
Age          11.2500     5.052      2.227      0.269     -52.939     75.439
-----
Omnibus:                     nan     Durbin-Watson:                  2.500
Prob(Omnibus):                  nan     Jarque-Bera (JB):            0.747
Skew:                          -0.913    Prob(JB):                      0.688
Kurtosis:                       2.500    Cond. No.                  6.38e+03
-----

```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 - [2] The condition number is large, 6.38e+03. This might indicate that there are strong multicollinearity or other numerical problems.
- -----

Interpretation:

In the abstract sense, the coefficient of x1 that is, the coefficient on education indicates that for every additional unit of education, the dependent variable y increases by approximately 15.8333 on avg

```
C:\Users\default\Desktop-GGCF6CQ\anaconda3\Lib\site-
packages\statsmodels\stats\stattools.py:74: ValueWarning: omni_normtest is not
valid with less than 8 observations; 5 samples were given.
  warn("omni_normtest is not valid with less than 8 observations; %i "%
```

```
[11]: print("""Answer to question #2""")
data = {
    'Dependent': [35, 50, 65, 70, 80],
```

```

'education': [12, 16, 18, 20, 21],
'experience': [5, 10, 12, 15, 18],
}

df = pd.DataFrame(data)
print(df)

Y=df ["Dependent"]
X=df [["education", "experience"]]
X=sm.add_constant(X)
model = sm.OLS(Y, X).fit()
print(model.summary())
print(""""
print("""
print(f"\nInterpretation:")
print("""When age is removed we get a coefficient of 2.976 on education when
it was previously 15.83, this leads to believe that age and education were
are positively correlated, and age was previously accounting for some of the
variation in the dependent variable that education now partially absorbs.
This illustrates the omitted variable bias, by showing that when you leave
relevant variables such as age it can change the effects of other variables
in your model such as with age here."""")
```

Answer to question #2:

	Dependent	education	experience
0	35	12	5
1	50	16	10
2	65	18	12
3	70	20	15
4	80	21	18

OLS Regression Results

```
=====
Dep. Variable: Dependent R-squared: 0.980
Model: OLS Adj. R-squared: 0.960
Method: Least Squares F-statistic: 49.34
Date: Sun, 02 Nov 2025 Prob (F-statistic): 0.0199
Time: 20:10:55 Log-Likelihood: -11.101
No. Observations: 5 AIC: 28.20
Df Residuals: 2 BIC: 27.03
Df Model: 2
Covariance Type: nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	-8.5204	28.760	-0.296	0.795	-132.263	115.222
education	2.9762	3.216	0.925	0.452	-10.863	16.816
experience	1.3946	2.325	0.600	0.610	-8.609	11.398

=====

Omnibus:	nan	Durbin-Watson:	3.570
Prob(Omnibus):	nan	Jarque-Bera (JB):	0.362
Skew:	-0.013	Prob(JB):	0.835
Kurtosis:	1.682	Cond. No.	401.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Interpretation:

When age is removed we get a coefficient of 2.976 on education when it was previously 15.83, this leads to believe that age and education were positively correlated, and age was previously accounting for some of the variation in the dependent variable that education now partially absorbs. This illustrates the omitted variable bias, by showing that when you leave relevant variables such as age it can change the effects of other variables in your model such as with age here.

```
C:\Users\default\Desktop-GGCF6CQ\anaconda3\Lib\site-
packages\statsmodels\stats\stattools.py:74: ValueWarning: omni_normtest is not
valid with less than 8 observations; 5 samples were given.
  warn("omni_normtest is not valid with less than 8 observations; %i "%
```

```
[12]: print("""Answer to question #3""")
data = {
    'Dependent': [35, 50, 65, 70, 80],
    'education': [12, 16, 18, 20, 21],
    'experience': [5, 10, 12, 15, 18],
    'Age': [25, 30, 32, 35, 40]
}

df = pd.DataFrame(data)
print(df)
Y=df["Dependent"]
X=df[["education", "experience", "Age"]]
X=sm.add_constant(X)
model = sm.OLS(Y, X).fit()
print("")-----#
print("")-----#
```

```
corr_matrix = df[['education', 'experience', 'Age']].corr()
print(corr_matrix)
```

```

print(""""
print("""
X = sm.add_constant(df[['education', 'experience', 'Age']])
vif = pd.DataFrame()
vif["Variable"] = X.columns
vif["VIF"] = [variance_inflation_factor(X.values, i)
              for i in range(X.shape[1])]
print(vif)
print("""
print("""
print(f"\nInterpretation:")
print("""The correlation matrix and VIF values reveal severe multicollinearity
    ↪among Education, Experience, and Age (all VIFs far exceed 10). This means
    ↪these variables contain highly overlapping information.

As a result, individual coefficient estimates are imprecise their standard
    ↪errors are inflated, and their magnitudes can vary widely when one regressor
    ↪is removed. This explains why the coefficient on Education changed
    ↪drastically when Age was excluded. This can also be seen from the
    ↪correlation matrix where each variable closely follows the other."""")

print("""
print("""
print("""Answer to question #4:""")
t_value = 2.5 / 0.8
p = 2 * (1-stats.t.cdf(abs(t_value), df=1))
alpha = 0.05
print(f"t_value = {t_value:.3f}, p = {p :.5f}")

if p < alpha:
    print("Reject H0: Education significantly affects y")
else:
    print("Fail to Reject H0: We dont have enough information to say that
        ↪Education significantly affects y")
print("""
print("""
print("""Answer to question #5:""")
b = 2.5
se = 0.8
df = 1
alpha = 0.05
t_crit = stats.t.ppf(1 - alpha/2, df)
lower = b - t_crit * se
upper = b + t_crit * se

print(f"95% Confidence Interval: ({lower:.3f}, {upper:.3f})")
print(f"\nInterpretation:")

```

```

print(f"Holding Experience and Age constant, we are 95% confident the true
    ↵effect of one more unit of Education lies between: ({lower:.3f}, {upper:.
    ↵3f}) Because 0 is inside the interval, Education is not statistically
    ↵significant at 5%.")

```

Answer to question #3:

	Dependent	education	experience	Age
0		35	12	5 25
1		50	16	10 30
2		65	18	12 32
3		70	20	15 35
4		80	21	18 40

	education	experience	Age
education	1.000000	0.988212	0.964229
experience	0.988212	1.000000	0.993065
Age	0.964229	0.993065	1.000000

	Variable	VIF
0	const	26125.000000
1	education	452.266667
2	experience	2298.916667
3	Age	766.850000

Interpretation:

The correlation matrix and VIF values reveal severe multicollinearity among Education, Experience, and Age (all VIFs far exceed 10). This means these variables contain highly overlapping information.

As a result, individual coefficient estimates are imprecise their standard errors are inflated, and their magnitudes can vary widely when one regressor is removed. This explains why the coefficient on Education changed drastically when Age was excluded. This can also be seen from the correlation matrix where each variable closely follows the other.

Answer to question #4:

t_value = 3.125, p = 0.19716

Fail to Reject H0: We dont have enough information to say that Education significantly affects y

Answer to question #5:

95% Confidence Interval: (-7.665, 12.665)

Interpretation:

Holding Experience and Age constant, we are 95% confident the true effect of one more unit of Education lies between: (-7.665, 12.665) Because 0 is inside the interval, Education is not statistically significant at 5%.